

Retail Insights Assistant

A Multi-Agent GenAI System for automated sales analytics, SQL generation, and business intelligence.

LangGraph Pipeline

4-Agent Orchestration with GPT-4

DuckDB Engine

High-performance local OLAP

Streamlit UI

Real-time status & visualization

Auto-Validation

Self-correcting data integrity

Agents



Query Resolution

NL → SQL Translation



Data Extraction

Executes queries & retrieves data



Validation Agent

Quality Assurance



Response Generation

Insight Narrative

Executive Summary



01_SYSTEM

What is it?

A Multi-agent GenAI pipeline designed for automated sales analytics.

- Flow: Query → Extract → Validate → Respond
- Converts natural language to SQL & executes on DuckDB



02_VALUE

Why it matters

Delivers reliable analytics where standard LLMs often hallucinate.

- ✓ Robust multi-level fallbacks & validation agents
- ✓ Business-friendly narratives with data limitations exposed



03_FEATURES

Key Features

Complete end-to-end solution from UI to Database.

- Streamlit UI with real-time agent status monitoring
- Smart visualization engine (Auto-chart selection)
- Cloud-ready & containerized deployment



04_IMPACT


Outcomes

Empowers stakeholders with instant data access.

- ↳ 10x faster ad-hoc analysis compared to manual SQL
- ↳ Consistent answers reducing operational risk

End-to-End Data Flow

● Flow Control Data Path ● AI Processing



LangGraph Orchestrator

orchestrator.py

</> Lines of Code: 166

Initialization & Workflow Graph

```
class RetailInsightsOrchestrator
```

```
StateGraph Compiled
```

```
API KEY
```

```
st.session_state
```

```
MODEL
```

```
gpt-4-turbo-preview
```

```
DATA PATH
```

```
Sales Dataset/
```

SEQUENTIAL STATEGRAPH PIPELINE

```
▶ Entry
```

```
Passing: AgentState  
{ query, sql, data... }
```

```
1
```

```
Query  
Resolution
```

```
2
```

```
Data  
Extraction
```

```
3
```

```
Validation
```

```
4
```

```
Response  
Gen
```

```
■ End
```

Agent System Overview



Input Processor



Executor

QueryResolutionAgent

Converts natural language into executable SQL using deep database context.

- Injects full schema & sample values into context window
- Handles DuckDB specifics (e.g., `YEAR()` vs `DATE_PART`)
- Generates complex YoY logic via self-joins

DataExtractionAgent

Executes queries and retrieves data with robust error handling.

- Safely executes SQL on local DuckDB engine
- **Smart Fallback:** If SQL fails, auto-fetches relevant summaries
- Attaches dataset metadata (date ranges, available quarters)



Quality Control



Narrator

ValidationAgent

Ensures data relevance and structural integrity before presentation.

- Structural checks for missing keys or empty payloads
- LLM-based semantic review: "Does this answer the question?"
- Assigns confidence score (0.0 - 1.0) to guide UI warnings

ResponseGenerationAgent

Crafts business-friendly insights and executive summaries.

- Highlights specific numbers, growth rates, and key trends
- Explains data limitations (e.g., "Q4 data is partial")
- Adapts tone for executive stakeholders (concise, direct)

Data Processing

DuckDB Engine

Engine & Schema Architecture



DuckDB OLAP Engine

Embedded analytics engine utilizing a persistent database file. Handles high-performance SQL queries directly on tabular data.

retail_data.duckdb

Auto-Load on Cold Start

1

amazon_sales

Amazon Sale Report.csv

Order ID (PK)

Date [DATE]

Amount [DOUBLE]

Category

Status

ship-state

2

inventory

Sale Report.csv

Category

Size

Color

Stock [INT]

3

international_sales

International sale Report.csv

CUSTOMER

Months

GROSS AMT

</> Core Methods & Cloud Logic

>_ Key Methods

Primary interface for Agent interaction

Execute raw SQL & return DataFrame

```
execute_query(sql: str) -> pd.DataFrame
```

Get schema + samples for LLM context

```
get_table_context(table_name: str)
```

Aggregate cross-table summary stats

```
get_summary_statistics()
```

Fallback data generators

```
get_top_categories(n)
```

```
getRegionalPerformance()
```



Cloud Deployment Strategy

Handling ephemeral file systems (Streamlit Cloud)

> Persistence Check

Checks information_schema.tables on init to avoid redundant re-loading.

Technology Stack

All dependencies pinned for reproducibility



Core Logic

Orchestration & LLM

- LangGraph v1.0.1
- LangChain v0.3.27
- OpenAI (SDK) v2.9.0
- LangChain-Core v0.3.80
- LC-OpenAI v0.3.35



Data & Viz

Processing & Analytics

- DuckDB v1.4.4
- Pandas v2.2.3
- NumPy v2.2.2
- Plotly v6.5.2



App & Utils

Frontend & Utilities

- Streamlit v1.49.1
- Pydantic v2.11.7
- tiktoken v0.9.0
- python-dotenv v1.0.1

Data Flow Processes


Q&A Logic

Summary Logic




Conversational Q&A Mode

Standard complex query pipeline with multi-agent orchestration



Executive Summary Mode

Pre-defined aggregation pipeline for instant dashboard generation



Deployment Options

Python 3.9+ Ephemeral DB



Local Dev

For development & testing

QUICK START



```
# Install dependencies
$ pip install -r requirements.txt

# Run Application
$ streamlit run app.py
```

CONFIGURATION

- ✓ Enter API key in sidebar UI
- ✓ Hot-reloading enabled
- ✓ DuckDB creates local file



Cloud Hosted

Streamlit Community Cloud

ZERO OPS



1. Push code to GitHub
2. Connect repo on share.streamlit.io
3. App deploys automatically

KEY FEATURES

- ✓ DuckDB rebuilds on cold start
- ✓ No secrets file management needed
- ✓ API Key entered by user at runtime



Docker Prod

Containerized Deployment

SCALABLE



```
FROM python:3.9-slim
WORKDIR /app
COPY . .
RUN pip install ...
EXPOSE 8501
CMD ["streamlit", "run", "app.py"]
```

PRODUCTION SPEC

- ✓ Ephemeral DuckDB filesystem
- ✓ Bundled CSVs for auto-load
- ✓ Ready for K8s / ECS

TARGET: 100GB+ DATA

Scalability Roadmap

Infrastructure Evolution
Single Node → Distributed Cluster



Data Layer

Embedded → Cloud Warehouse

MIGRATE

* Snowflake / BigQuery

Year Partitioning

Parquet

Processing

Pandas/Single

Distributed ETL

SCALE

⚡ PySpark / Dask

Apache Airflow

Infrastructure

Streamlit Cloud

Container Cluster

DEPLOY

Docker

K8s HPA

Redis Cache

LLM Optimization

Direct API

Smart Routing

OPTIMIZE

Model Routing

3.5 GPT-3.5 vs 4

Semantic Cache

Retrieval

SQL Only

Hybrid Search

ENHANCE

RAG

FAISS / Pinecone

Summaries

Monitoring

Basic Logging

Observability

TRACK

Prometheus

Grafana

Alerts